

Linux System Call

Practical Class 5

Systems and Storage Laboratory
Department of Computer Science and Engineering
Chung-Ang University

Index

❖ Creating a new system call

1. Download Kernel
2. Define system call function
3. Register system call number to the table
4. Register system call to the header file
5. Kernel compile & change
6. System call usage

Download Kernel (Practical class 4)

❖ Check current kernel version

- `$ uname -r`

❖ Download kernel code from <https://www.kernel.org/>

❖ Or use following command:

```
$ wget https://mirrors.edge.kernel.org/pub/linux/kernel/v5.x/linux-5.x.tar.gz  
$ tar -xvf linux-5.x.tar.gz
```

Define system call function

❖ Create system call function

- `$ vim linux-x.x.x/kernel/my_syscall.c` (Make new file)

```
1 #include <linux/syscalls.h>
2
3 SYSCALL_DEFINE0(mycall)
4 {
5     printk("System Call Example!\n");
6
7     return 0;
8 }
```


❖ Add this code to Makefile

- `$ vim linux-x.x.x/kernel/Makefile`

```
obj-y = fork.o exec_domain.o panic.o \
      cpu.o exit.o softirq.o resource.o \
      sysctl.o sysctl_binary.o capability.o ptrace.o user.o \
      signal.o sys.o umh.o workqueue.o pid.o task_work.o \
      extable.o params.o \
      kthread.o sys_ni.o nsproxy.o \
      notifier.o ksysfs.o cred.o reboot.o \
      async.o range.o smptboot.o ucount.o my_syscall.o
```

Register system call number to the table

- ❖ We need to add the system call number to the table.
- ❖ There are 2 popular architecture:
 - X86: Intel Processor
 - Arm64: Apple M1, M2, some AMD, smart tablet processors
- ❖ The system call table is in different files based on the architecture. You should check the processor architecture before implement the next step.




Register system call number to the table

❖ x86

- `$ vim linux-x.x.x/arch/x86/entry/syscalls/syscall_64.tbl`

```
...  
<NUM>      common  mycall      __x64_sys_mycall  # add syscall number  
...
```

```
424  common  pidfd_send_signal  __x64_sys_pidfd_send_signal  
425  common  io_uring_setup      __x64_sys_io_uring_setup  
426  common  io_uring_enter       __x64_sys_io_uring_enter  
427  common  io_uring_register    __x64_sys_io_uring_register  
428  common  open_tree            __x64_sys_open_tree  
429  common  move_mount          __x64_sys_move_mount  
430  common  fsopen              __x64_sys_fsopen  
431  common  fsconfig            __x64_sys_fsconfig  
432  common  fsmount             __x64_sys_fsmount  
433  common  fspick              __x64_sys_fspick  
434  common  pidfd_open          __x64_sys_pidfd_open  
435  common  clone3              __x64_sys_clone3/ptregs
```



```
424  common  pidfd_send_signal  __x64_sys_pidfd_send_signal  
425  common  io_uring_setup      __x64_sys_io_uring_setup  
426  common  io_uring_enter       __x64_sys_io_uring_enter  
427  common  io_uring_register    __x64_sys_io_uring_register  
428  common  open_tree            __x64_sys_open_tree  
429  common  move_mount          __x64_sys_move_mount  
430  common  fsopen              __x64_sys_fsopen  
431  common  fsconfig            __x64_sys_fsconfig  
432  common  fsmount             __x64_sys_fsmount  
433  common  fspick              __x64_sys_fspick  
434  common  pidfd_open          __x64_sys_pidfd_open  
435  common  clone3              __x64_sys_clone3/ptregs  
436  common  mycall              __x64_sys_mycall
```

Before

After

Register system call number to the table

❖ Arm64

- `$ vim linux-x.x.x/include/uapi/asm-generic/unistd.h`

❖ Use 460 for the mycall number (if you have error, increase higher)


❖ `__NR_syscalls` should be 461 (mycall +1)

```
#ifdef __ARCH_WANT_SYS_CLONE3
#define __NR_clone3 435
__SYSCALL(__NR_clone3, sys_clone3)
#endif

#undef __NR_syscalls
#define __NR_syscalls 436

/*
 * 32 bit systems traditionally used different
 * syscalls for off_t and loff_t arguments, whi
```

Before



```
#ifdef __ARCH_WANT_SYS_CLONE3
#define __NR_clone3 435
__SYSCALL(__NR_clone3, sys_clone3)
#endif

#define __NR_mycall 460
__SYSCALL(__NR_mycall, sys_mycall)

#undef __NR_syscalls
#define __NR_syscalls 461
```

After

Register system call to header file

- ❖ Add our system call to syscalls.h header file
 - `$ vim linux-x.x.x/include/linux/syscalls.h`
 - Add our system call function

```
...  
asmlinkage long sys_mycall(void);  
#endif /* CONFIG_ARCH_HAS_SYSCALL_WRAPPER */
```

```
1219 /*  
1220  * Not a real system call, but a placeholder for syscalls which are  
1221  * not implemented -- see kernel/sys_ni.c  
1222  */  
1223 asmlinkage long sys_ni_syscall(void);  
1224  
1225 asmlinkage long sys_mycall(void);  
1226  
1227 #endif /* CONFIG_ARCH_HAS_SYSCALL_WRAPPER */
```


Kernel compile

❖ Kernel compile

- `$ make menuconfig` (if it is not configured)
- `$ make -j 4`
- `$ make modules -j 4`
- `$ sudo make modules_install -j 4`
- `$ sudo make install`

Change the default boot option

❖ *If the grub is not configured for the new compiled kernel, you can follow the below steps:*

❖ **Check grub id from there**

- `$ vim /boot/grub/grub.cfg`

❖ **Change grub_default setting**

- `$ vim /etc/default/grub`

```
GRUB_DEFAULT="submenu_id>menuentry_id"
```

Ex)

```
GRUB_DEFAULT="gnulinux-advanced-65c9af03-3d9b-411c-99b2-a9ada0961a40>gnulinux-4.7.0-1-amd64-advanced-65c9af03-3d9b-411c-99b2-a9ada0961a4"
```

❖ **Apply grub configuration**

- `$ update-grub`

System call usage

❖ sys_mycall usage example

- `$ vim syscall_example.c`
 - For the x86 version, the syscall number is 436
For the arm64 version, the syscall number is 460

```
1 #include <stdio.h>
2 #include <sys/syscall.h>
3
4 int main(void)
5 {
6     long ret = syscall(436);
7     printf("System Call returned: %ld\n", ret);
8
9     return 0;
10 }
```

- `$ gcc syscall_example.c -o result`

System call usage

❖ my_syscall usage example

- `$./result`

```
syslab@syslab-VirtualBox: ~  
syslab@syslab-VirtualBox:~$ gcc syscall_example.c -o result  
syscall_example.c: In function 'main':  
syscall_example.c:6:13: warning: implicit declaration of function  
'syscall' [-Wimplicit-function-declaration]  
    6 |     long ret = syscall(436);  
      |                   ^~~~~~  
syslab@syslab-VirtualBox:~$ ./result  
System Call returned: 0  
syslab@syslab-VirtualBox:~$ _
```

- `$ sudo dmesg`

```
[ 7.019395] *** VALIDATE vboxsf ***  
[ 7.019399] vboxsf: Successfully loaded version 6.1.38 r153438  
[ 7.019443] vboxsf: Successfully loaded version 6.1.38 r153438 on 5.4.214syslab SMP  
[ 7.021038] 13:07:04.467452 automount vbsvcAutomounterMountIt: Successfully mounted  
[ 12.393764] rfkill: input handler disabled  
[ 143.796105] hrtimer: interrupt took 10988967 ns  
[ 166.464380] System Call Example!
```