

Linux kernel data structure XArray

Practical Class 8

**Systems and Storage Laboratory
Department of Computer Science and Engineering
Chung-Ang University**

Index

❖ Introducing XArray

❖ XArray data structure

- XArray structure
- XArray node
- XArray node slots
- XArray node insertion routines
- XArray Normal APIs

❖ Xarray Usage Example

- ds_monitoring

Introducing XArray

❖ eXtensible Arrays

- Abstract data type which behaves like a very large array of pointers
- Normal pointers may be stored in the XArray directly
- They must be 4-byte aligned, which is true for any pointer returned from `kmalloc()` and `alloc_page()`
- Meets many of the same needs as a hash or a conventional resizable array
- More memory-efficient, parallelisable and cache friendly than a doubly-linked list
- Uses RCU and an internal spinlock to synchronize access
- The most important user of the XArray is the page cache

XArray structure

❖ The anchor of the XArray

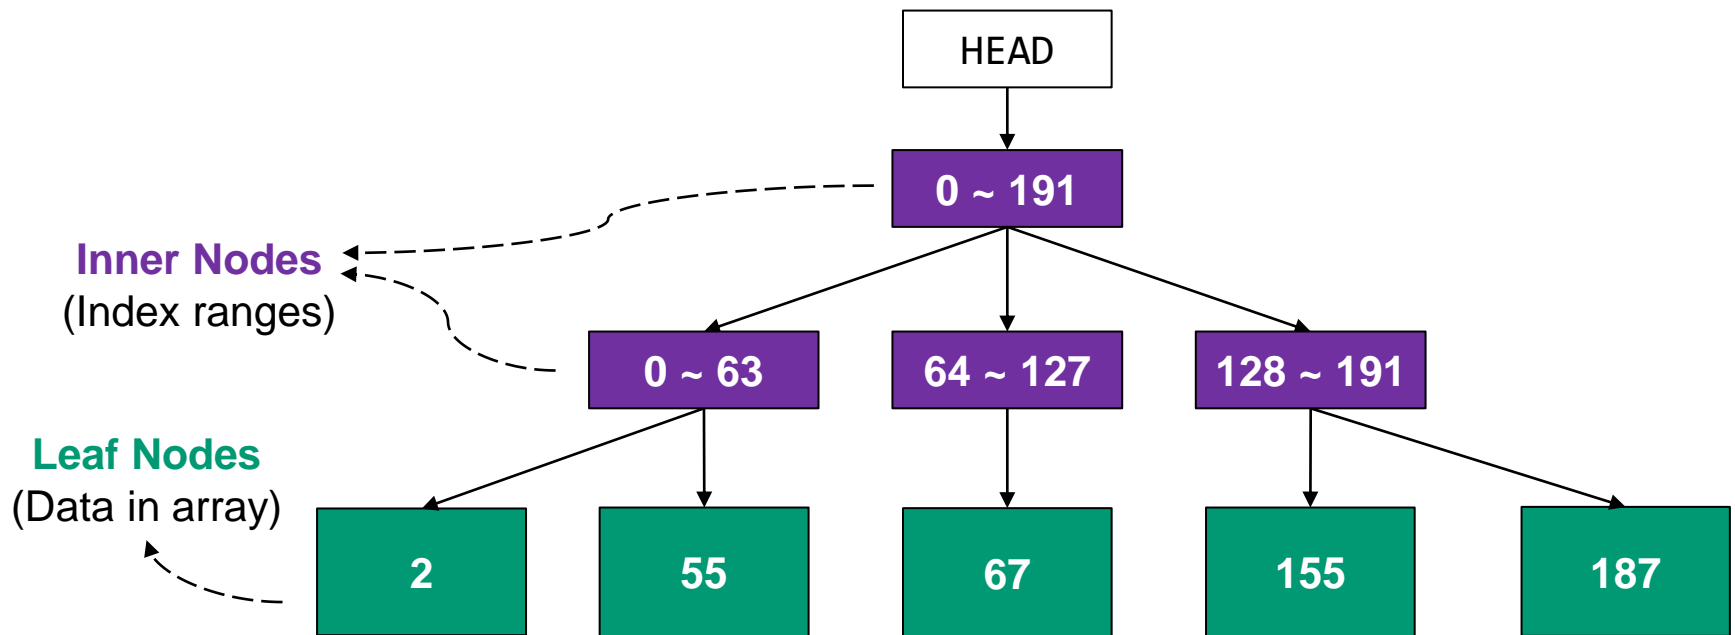
- It is a very small data structure so it could be easily defined statically or embed

```
struct xarray {  
    spinlock_t  xa_lock;  
    gfp_t       xa_flags;  
    void __rcu * xa_head;  
};
```

- **xa_lock**: could be used to protect the contents of XArray
- **xa_head**: the root node of XArray

XArray node

- ❖ The data structure of XArray nodes is **struct xa_node**
 - The data is only kept in leaf nodes
 - Inner nodes save index ranges
 - Inner nodes are only for navigating



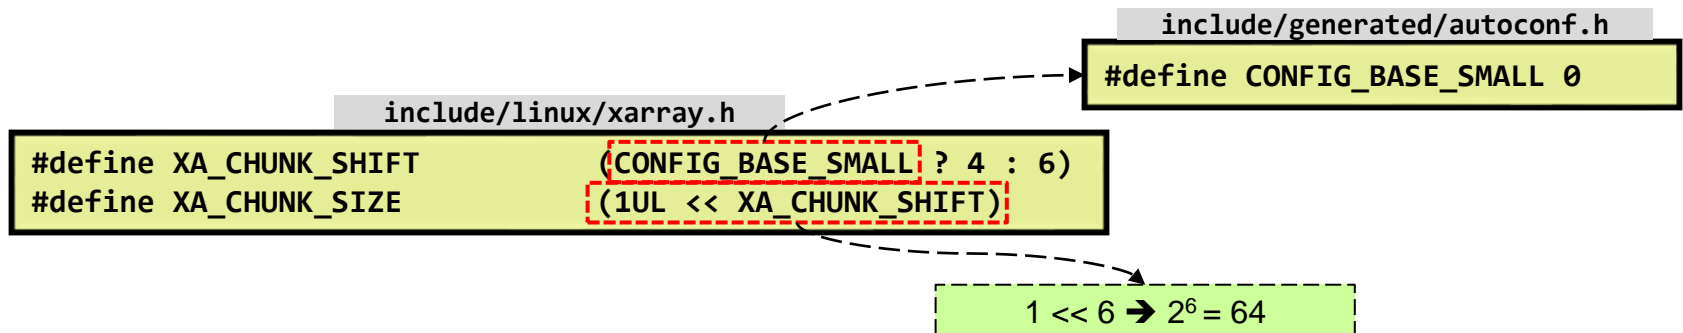
struct xa_node

❖ An element in XArray

```
struct xa_node {
    unsigned char    shift;           /* Bits remaining in each slot */
    unsigned char    offset;          /* Slot offset in parent */
    unsigned char    count;           /* Total entry count */
    unsigned char    nr_values;       /* Value entry count */
    struct xa_node __rcu *parent;      /* NULL at top of tree */
    struct xarray    *array;          /* The array we belong to */
    union {
        struct list_head private_list; /* For tree user */
        struct rcu_head rcu_head;      /* Used when freeing node */
    };
    void __rcu    *slots[XA_CHUNK_SIZE];
    union {
        unsigned long    tags[XA_MAX_MARKS][XA_MARK_LONGS];
        unsigned long    marks[XA_MAX_MARKS][XA_MARK_LONGS];
    };
};
```

struct xa_node


- **shift:** bits remaining in each slot (4 or 6)
- **offset:** the slot offset in **parent**
- **count:** the count of every non-NULL element in the slots array
- **nr_values:** the count of every element in slots which is either a value entry or a sibling of a value entry
- **array:** the xarray that the nodes belong to
- **slots:** an array saving children nodes with XA_CHUNK_SIZE elements
 - Default of XA_CHUNK_SIZE is 64 – with XA_CHUNK_SHIFT = 6



struct xa_node

❖ 64bit Architecture

- 576bytes per one xa_node structure
- 4KB page can have up to 7 nodes




struct xa_node

XArray node slots

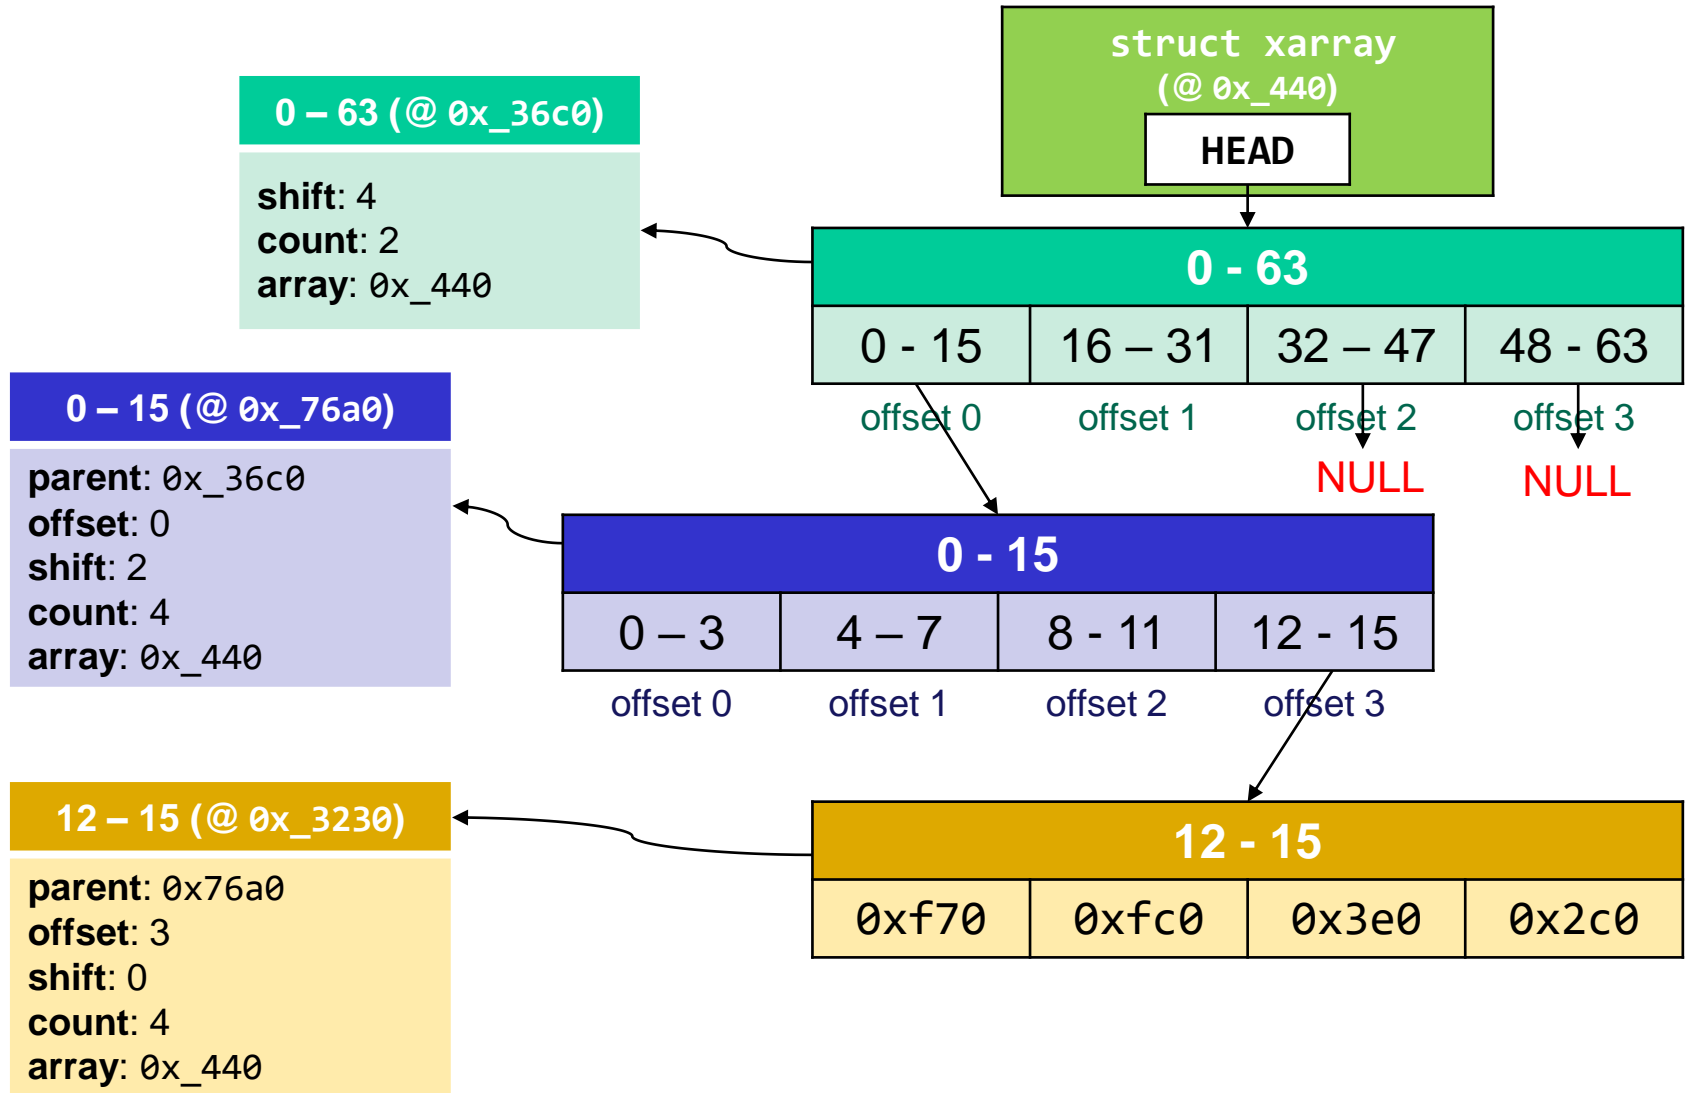
❖ Slots can contain inner nodes or leaf nodes

- Ex) Let $\text{XA_CHUNK_SIZE} = 4$
 - One node has 4 slots
 - $\text{XA_CHUNK_SHIFT} = 2$

```
#define XA_CHUNK_SIZE      (1UL << XA_CHUNK_SHIFT)
```




XArray node slots



XArray node insertion routines

❖ Initializing Xarray



```
struct xarray {  
    spinlock_t  xa_lock;  
    gfp_t       xa_flags;  
    void __rcu * xa_head;  
};
```

Xarray structure


xa_head	NULL
xa_lock	initialized
xa_head	NULL

Xarray node structure fields

shift	0
offset	0
count	0
nr_values	0
parent	NULL
xarray	NULL

XArray node insertion routines

❖ Insert first element




Xarray structure	
xa_head	NULL
xa_lock	initialized
xa_head	page0 address

```
struct xarray {  
    spinlock_t xa_lock;  
    gfp_t      xa_flags;  
    void __rcu * xa_head;  
};
```

Page address

XArray node insertion routines

❖ Insert second element



Xarray structure


xa_head	NULL
xa_lock	initialized
xa_head	node0 address

Xarray node structure fields

shift	0	
offset	0	
count	2	
nr_values	0	
parent	head	
array	0x01	
slots	[0]	P0 addr
	[1]	P1 addr

XArray node insertion routines

❖ Insert 64th element



Xarray structure

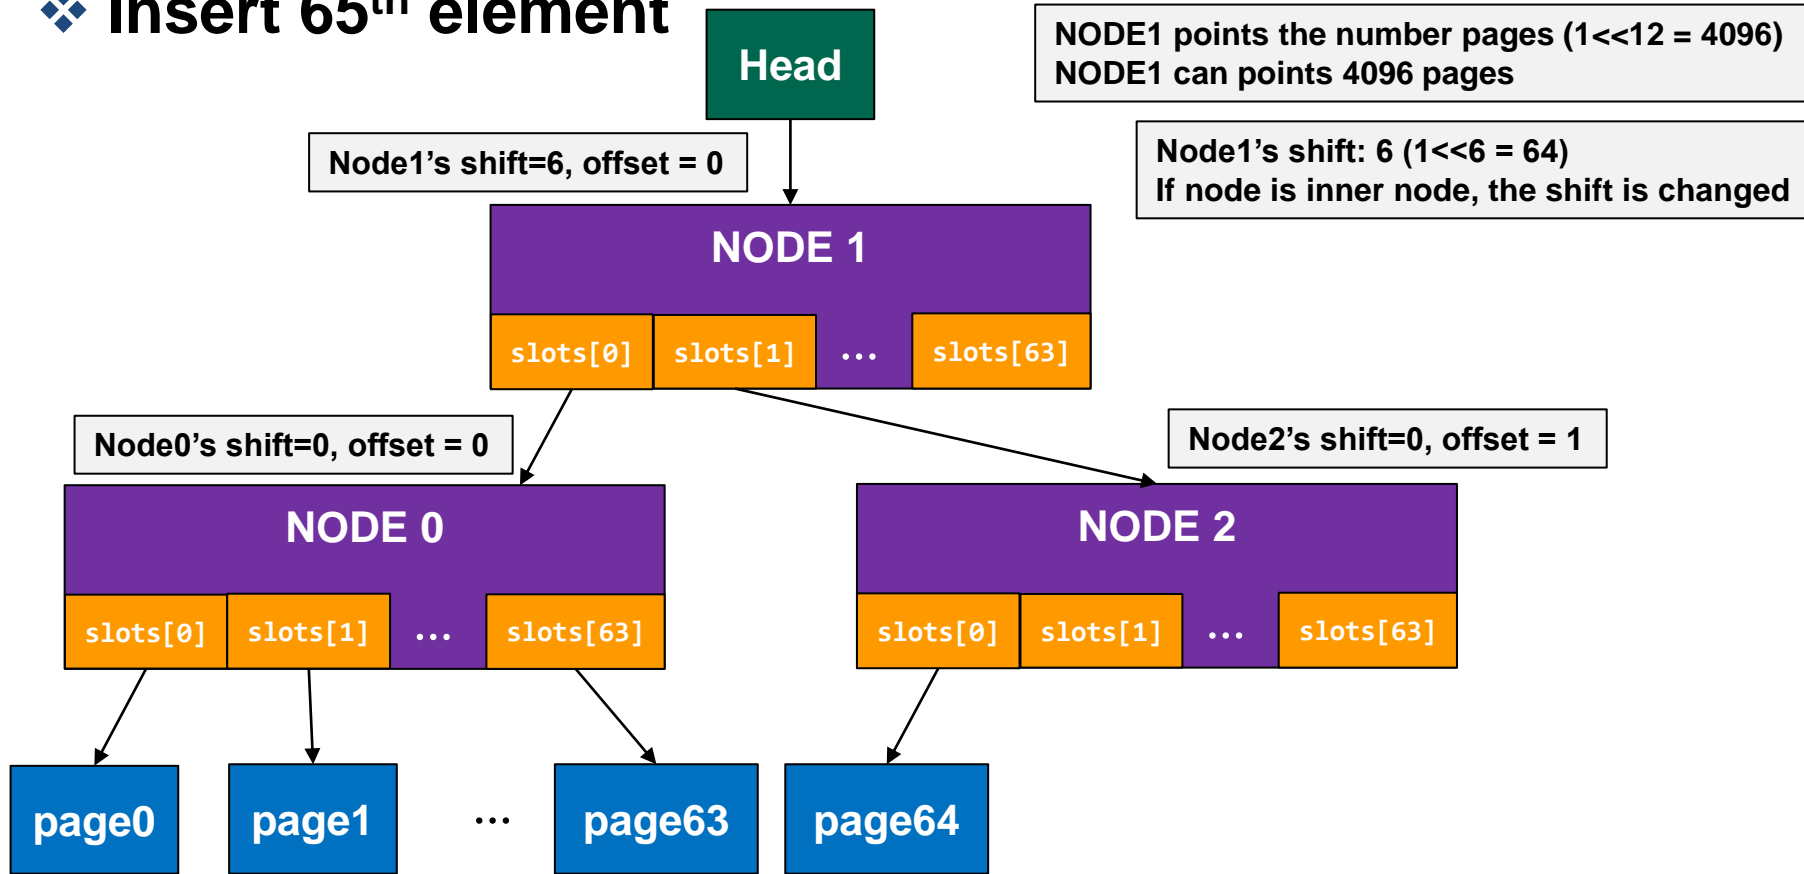
xa_head	NULL
xa_lock	initialized
xa_head	node0 address

Xarray node structure fields

shift	0
offset	0
count	64
nr_values	0
parent	head
array	0x01
slots[0]~[63]	P0 addr ~ P63 addr

XArray node insertion routines

❖ Insert 65th element



Xarray structure: `xa_head = NULL`; `xa_lock` is initialized; `xa_head = page address`
Xarray node structure 0 fields: `shift=0, offset=0, count=64, nr_values=0, parent=node1, xarray = 0x01, slots[0] ~ slots[63] = P0 addr ~ P63 addr`
Xarray node structure 1 fields: `shift=6, offset=0, count=2, nr_values=0, parent=head, xarray = 0x01, slots[0] = node0, slots[1] = node2`
Xarray node structure 2 fields: `shift=0, offset=1, count=1, nr_values=0, parent=node1, xarray = 0x01, slots[0] = P64 addr`

XArray Normal APIs

❖ Search

- Search the XArray for an entry

```
void *xa_find(struct xarray *xa, unsigned long *indexp,  
             unsigned long max, xa_mark_t filter)
```

- **xa**: XArray to search entry from
 - **indexp**: Pointer to an index. (Found entry index will be saved here)
 - **max**: Maximum index to search to.
 - **filter**: Selection criterion.
 - Return: The entry, if found, otherwise **NULL**
- Load an entry from an XArray

```
void *xa_load(struct xarray *xa, unsigned long index);
```

- **xa**: XArray to load entry from
- **index**: Entry index that we are searching for
- Return: The entry at **index** in **xa**.

XArray Normal APIs

❖ Insert

- Store this entry in the XArray unless another entry is already present

```
int xa_insert(struct xarray *xa, unsigned long index,  
             void *entry, gfp_t gfp)
```

- **xa**: XArray.
- **index**: Index into array.
- **entry**: New entry.
- **gfp**: Memory allocation flags
- Return:
 - ✓ **0** if the store succeeded.
 - ✓ **-EBUSY** if another entry was present.
 - ✓ **-ENOMEM** if memory could not be allocated.

XArray Normal APIs

❖ Update

- Store this entry in the XArray

```
void *xa_store(struct xarray *xa, unsigned long index,  
              void *entry, gfp_t gfp)
```

- **xa**: XArray.
- **index**: Index into array to update with.
- **entry**: New entry to store.
- **gfp**: Memory allocation flags.
- Return:
 - ✓ The old entry at this index on success
 - ✓ **xa_err(-EINVAL)** if **entry** cannot be stored in an XArray
 - ✓ **xa_err(-ENOMEM)** if memory allocation failed

XArray Normal APIs

❖ Delete

- Erase this entry from the Xarray

```
void *xa_erase(struct xarray *xa, unsigned long index);
```

- **xa**: XArray.
- **index**: Index of entry to delete from **xa**.
- **Return**: The entry which used to be at this index

ds_monitoring

Xarray Usage Example

Purpose of ds_monitoring

❖ To monitor values in overall data structure

- Functions in the Linux kernel are called too many times
- There is a limit to checking values only with `printk()`
- Collecting the data in one place and printing them out all at once later will be one solution

❖ XArray will easily handle the problem, according to the goal

- XArray is Linux kernel's native data structure
- It has advantages in terms of stability, reliability and efficiency.

struct ds_monitoring structure

❖ Main container for the XArray of elements

```
struct ds_monitoring {  
    struct xarray *elements;  
    unsigned long long total_counts;  
    const struct ds_monitoring_operations *dm_ops;  
};
```

- **elements**: pointer to the XArray which contains pointers to each ds_monitoring_elem
- **total_counts**: counter for how many times have xarray been called
- **dm_ops**: member functions for manipulating and managing elements

struct ds_monitoring_operations

❖ Collection of pointers to functions for manipulating and managing elements

- User must implement functions which fit for purpose, and connect them to this structure
- Each function pointers will be embedded into ds_monitoring's own functions

```
struct ds_monitoring_operations {  
    unsigned long (*get_index)(void *elem);  
    const char * (*get_name)(void *elem);  
    void (*print_elem)(  
        unsigned long index,  
        const char *name,  
        unsigned long long count,  
        int percentage  
    );  
};
```

struct ds_monitoring_elem

❖ Element of XArray

```
struct ds_monitoring_elem {  
    unsigned long key;  
    char *name;  
    unsigned long long count;  
};
```

- **key**: index of the element itself in XArray
- **name**: representative name for this element, will be used when printing elements
- **count**: counter for how many times have this element been searched from XArray

DEFINE_DS_MONITORING

❖ Defines a Data Structure Monitoring

- This is intended for file scope definitions of Data Structure Monitoring.
- It declares and initializes an empty Data Structure Monitoring structure with the chosen name
- It does the initialization at compile time instead of runtime

```
#define DEFINE_DS_MONITORING(name, get_idx_fn, get_name_fn, print_fn) \
    DEFINE_XARRAY(name##_xarray); \
    DEFINE_DS_MONITORING_OPS(name, get_idx_fn, get_name_fn, print_fn); \
    struct ds_monitoring name = DS_MONITORING_INIT(name##_xarray, name##_dm_ops);
```

- **name**: A string that names your Data Structure Monitoring
- **get_idx_fn**: Address of function which prepares key for a ds_monitoring_elem
- **(Optional) get_name_fn**: Address of function which determines name of each ds_monitoring_elem (pass **NULL** if not used)
- **print_fn**: A function address which prints out for each element

DEFINE_DS_MONITORING_OPS

- ❖ Defines `ds_monitoring_operations` and map each function pointers with real function addresses

```
#define DEFINE_DS_MONITORING_OPS(name, get_idx_fn, get_name_fn, print_fn) \
    static const struct ds_monitoring_operations name##_dm_ops = { \
        .get_index = get_idx_fn, \
        .get_name = get_name_fn, \
        .print_elem = print_fn, \
    }
```

- **name**: A string of your `ds_monitoring` name
- **get_idx_fn**: A function address which prepares key for a `ds_monitoring_elem`
- **get_name_fn**: A function address which determines name of each `ds_monitoring_elem`
- **print_fn**: A function address which prints out for each element

DS_MONITORING_INIT

❖ Macro for Initializing ds_monitoring

```
#define DS_MONITORING_INIT(xarray, _dm_ops) \
{ \
    .elements = &xarray \
    .total_counts = 0, \
    .dm_ops = &_dm_ops, \
}
```

- **xarray**: XArray to map with elements in ds_monitoring
- **_dm_ops**: ds_monitoring_operations to map with dm_ops in ds_monitoring

DECLARE_DS_MONITORING

- ❖ Declares ds_monitoring which was previously defined in another file

```
#define DECLARE_DS_MONITORING (name) \
    extern struct ds_monitoring name;
```

find_ds_monitoring()

❖ Find if ds_monitoring_elem of elem is in Xarray

- If found, increase referenced count
- If not, create a new ds_monitoring_elem for elem

```
void find_ds_monitoring(struct ds_monitoring *dm, void *elem)
{
    struct ds_monitoring_elem *cur;
    unsigned long xa_index;

    if (dm->dm_ops->get_index) {
        xa_index = dm->dm_ops->get_index(elem);

        // search ds_monitoring_elem at index of xa_index from xarray root
        cur = (struct ds_monitoring_elem *) xa_load(dm->elements, xa_index);

        if (cur) {
            __sync_fetch_and_add(&cur->count, 1);
        } else {
            insert_ds_monitoring(dm, xa_index, elem);
        }
        __sync_fetch_and_add(&dm->total_counts, 1);
    }
}
```

insert_ds_monitoring()

- ❖ Creates new xa_mapping_elem from elem and append it into the XArray

```
static void
insert_ds_monitoring(struct ds_monitoring *dm, unsigned long index, void *elem)
{
    char *name;
    struct ds_monitoring_elem *new = kmalloc(sizeof(struct ds_monitoring_elem),
                                              GFP_KERNEL);

    new->key = index;
    new->count = 1;

    if (dm->dm_ops->get_name) {
        name = dm->dm_ops->get_name(elem);
        new->name = kmalloc(strlen(name)+1, GFP_KERNEL);
        strcpy(new->name, name);
    } else {
        new->name = NULL;
    }
    xa_store(dm->elements, new->key, (void*) new, GFP_KERNEL);
}
```

print_ds_monitoring()

- ❖ Dump all the ds_monitoring_elem information in the ds_monitoring

```
void print_ds_monitoring(struct ds_monitoring* dm)
{
    unsigned long cur_idx;
    void *cur;
    char *cur_name;
    unsigned long long cur_count;
    int percentage;

    if (!dm->total_counts)
        return;

    xa_for_each(dm->elements, cur_idx, cur) {
        cur_name = ((struct ds_monitoring_elem *)cur)->name;
        cur_count = ((struct ds_monitoring_elem *)cur)->count;
        percentage = cur_count * 100 / dm->total_counts;
        dm->dm_ops->print_elem(cur_idx, cur_name, cur_count, percentage);
    }
}
```

delete_ds_monitoring()

- ❖ Frees all the allocated memory for ds_monitoring, and then destroys it

```
void delete_ds_monitoring(struct ds_monitoring *dm)
{
    unsigned long cur_idx;
    void *cur;
    char *cur_name;

    xa_for_each(dm->elements, cur_idx, cur) {
        cur_name = ((struct ds_monitoring_elem *)cur)->name;
        kfree(cur_name);
        kfree(cur);
    }
    xa_destroy(dm->elements);
}
```


Usage Example: Definition

- ❖ Define your own `ds_monitoring` at the **global scope** of the `c` file you want to use in
 - In this case, struct `ds_monitoring thread_dm` will be created

```
DEFINE_DS_MONITORING(thread_dm, get_thread_idx, get_thread_name, print_zone_dm);
```

Usage Example: get_idx_fn

❖ get_idx_fn should

- Receive void* of any structure that you want to watch
- Return unsigned long value of index of which element should become

```
static unsigned long get_thread_idx(void *elem)
{
    int node_idx;
    unsigned long xa_index;
    int zone_idx;

    struct task_struct *current_task = (struct task_struct *) elem;

    return (unsigned long) current_task->pid;
}
```

Usage Example: get_name_fn

❖ get_name_fn should

- Receive void* of any structure that you want to watch
- Return const char* value of name by which element should be called

```
static const char * get_thread_name(void *elem)
{
    struct task_struct *current_task = (struct task_struct *) elem;
    return current_task->comm;
}
```

Usage Example: print_fn

❖ print_fn should receive

- unsigned long **index**: key value for each element
- const char ***name**: representative name for this element
- unsigned long long **count**: element has been searched
- int **percentage**: proportion of this element count in the total_counts

```
static void print_zone_dm(unsigned long pid,
    const char *name,
    unsigned long long count,
    int percentage)
{
    printk("thread %s: pid %ld called wb_check_background_flush() \
        %lld times (%d%%)\n", name, pid, count, percentage);
}
```

Usage Example: find_ds_monitoring

❖ User will try to find some element in the ds_monitoring

```
#include "ds_monitoring.h"

DEFINE_DS_MONITORING(thread_dm, get_thread_idx, get_thread_name, print_zone_dm);

static long wb_check_background_flush(struct bdi_writeback *wb)
{
    if (wb_over_bg_thresh(wb)) {
        struct wb_writeback_work work = {
            .nr_pages    = LONG_MAX,
            .sync_mode    = WB_SYNC_NONE,
            .for_background = 1,
            .range_cyclic  = 1,
            .reason        = WB_REASON_BACKGROUND,
        };

        find_ds_monitoring(&thread_dm, current);

        return wb_writeback(wb, &work);
    }
    return 0;
}
```

Usage Example: print dm & delete dm

- ❖ Later at some point, you can print out the elements in the ds_monitoring
- ❖ After finishing utilizing ds_monitoring, make sure you destroy it

```
#include "ds_monitoring.h"

DECLARE_DS_MONITORING(thread_dm);

static void __exit ext4_exit_fs(void)
{
    ...

    print_ds_monitoring(&thread_dm);
    delete_ds_monitoring(&thread_dm);
}
```

Usage Example: Dump result

```
thread kworker/u131:2: pid 61687 called wb_check_background_flush() 1 times (10%)
thread kworker/u130:2: pid 62514 called wb_check_background_flush() 1 times (10%)
thread kworker/u129:2: pid 64757 called wb_check_background_flush() 1 times (10%)
thread fio: pid 69653 called wb_check_background_flush() 2 times (20%)
thread fio: pid 69654 called wb_check_background_flush() 2 times (20%)
thread fio: pid 69655 called wb_check_background_flush() 2 times (20%)
thread fio: pid 69656 called wb_check_background_flush() 1 times (10%)
```

Overall Form

❖ Monitoring part

```
#include "ds_monitoring.h"

static unsigned long get_target_idx(void *elem)
{
    ...
}

static const char * get_target_name(void *elem)
{
    ...
}

static void print_target_dm(unsigned long index, const char *name,
                           unsigned long long count, int percentage)
{
    ...
}

DEFINE_DS_MONITORING(your_dm, get_target_idx, get_target_name, print_target_dm);

int function(...)
{
    struct data_structure target_to_monitor = {
        .foo    = LONG_MAX,
        .bar    = "Hello",
    };
    ...
    find_ds_monitoring(&your_dm, &target_to_monitor);
    ...
    return 0;
}
```


Overall Form

❖ Dumping result and destroying part

```
#include "ds_monitoring.h"

DECLARE_DS_MONITORING(your_dm);

static void __exit module_exit_fs(void)
{
    ...

    print_ds_monitoring(&your_dm);

    delete_ds_monitoring(&your_dm);
}
```