

# Assignment-8

## Linux System and its Applications

Systems and Storage Laboratory

Department of Computer Science and Engineering

Chung-Ang University

# Assignment-8: Checking CPU utilization

- ❖ **First, printk() the CPU id in pxt4\_file\_write\_iter()**
  1. Build pxt4 module and mount the testing device with pxt4
    - **Free space** in the device must be larger than **3 \* (RAM size)**
  2. Run Fio test with the following condition (next page)
    - Refer to [Practical Class 7-a]
  3. Unmount the device and remove pxt4 module
  4. Use `dmesg` command to check the result
  
- ❖ **Then use ds\_monitoring to check how many times pxt4\_file\_write\_iter() have been called from each CPU**
  - Do the same above 1 ~ 4

# Assignment-8: Checking CPU utilization

## ❖ Test configuration should satisfy the below conditions

- Buffered Sequential Write
- Block Size = 4K
- Numjobs = machine CPU cores
- Total size = 3 times larger than your memory size

### ■ Example Fio Script

```
; -- start job file --  
[global]  
name=<job name>  
directory=<pvt4 filesystem mount point (e.g. /mnt/test)>  
rw=write  
bs=4K  
direct=0  
numjobs=<number of CPU cores>  
verify=meta  
  
[fio-test]  
size=<(RAM size) * 3 / (numjobs)>  
group_reporting  
; -- end job file --
```

# Assignment-8: Checking CPU utilization

## ❖ What to handout

- Take a screenshot of
  1. First dmesg result from printk
    - ✓ You don't have to include whole lines from pxt4.
    - ✓ Just last lines are okay.
  2. The ds\_monitoring result that contains
    - ✓ CPU id
    - ✓ Function call counts
    - ✓ Overall function call rate percentage

## ❖ Submit within pdf format

- Make sure to include your name and student id

# Example screenshot 1

## ❖ A lot of printk()s in dmesg, difficult to track

```
[22305.031685] cpu[3] called pxt4_file_write_iter()
[22305.031688] cpu[3] called pxt4_file_write_iter()
[22305.031692] cpu[3] called pxt4_file_write_iter()
[22305.031695] cpu[3] called pxt4_file_write_iter()
[22305.031699] cpu[3] called pxt4_file_write_iter()
[22305.031703] cpu[3] called pxt4_file_write_iter()
[22305.031706] cpu[3] called pxt4_file_write_iter()
[22305.031710] cpu[3] called pxt4_file_write_iter()
[22305.031713] cpu[3] called pxt4_file_write_iter()
[22305.031717] cpu[3] called pxt4_file_write_iter()
[22305.031720] cpu[3] called pxt4_file_write_iter()
[22305.031728] cpu[3] called pxt4_file_write_iter()
[22305.031733] cpu[3] called pxt4_file_write_iter()
[22305.031737] cpu[3] called pxt4_file_write_iter()
[22305.031740] cpu[3] called pxt4_file_write_iter()
[22305.031744] cpu[3] called pxt4_file_write_iter()
[22305.031748] cpu[3] called pxt4_file_write_iter()
[22305.031751] cpu[3] called pxt4_file_write_iter()
[22305.031754] cpu[3] called pxt4_file_write_iter()
[22305.031758] cpu[3] called pxt4_file_write_iter()
[22305.031761] cpu[3] called pxt4_file_write_iter()
[22305.031765] cpu[3] called pxt4_file_write_iter()
[22305.031768] cpu[3] called pxt4_file_write_iter()
[22305.031772] cpu[3] called pxt4_file_write_iter()
[22305.031776] cpu[3] called pxt4_file_write_iter()
[22305.031780] cpu[3] called pxt4_file_write_iter()
[22305.031783] cpu[3] called pxt4_file_write_iter()
[22305.031787] cpu[3] called pxt4_file_write_iter()
[22305.031790] cpu[3] called pxt4_file_write_iter()
[22305.031795] cpu[3] called pxt4_file_write_iter()
[22305.031799] cpu[3] called pxt4_file_write_iter()
[22305.031802] cpu[3] called pxt4_file_write_iter()
[22518.570654] file_write_iter is called 3,145,728 times, and the time interval is 172,919,752,964ns
syslab@syslab-VirtualBox:~/pxt4$
```

# Example screenshot 2

## ❖ Using ds\_monitoring

```
[ 12.615914] audit: type=1400 audit(1667452613.504:10): apparmor="STATUS" operation="profile_load" profile="unconfined" name="/usr/bin/man" pid=565 comm="apparmor_parser"
[ 12.615917] audit: type=1400 audit(1667452613.504:11): apparmor="STATUS" operation="profile_load" profile="unconfined" name="man_filter" pid=565 comm="apparmor_parser"
[ 13.772874] e1000: enp0s3 NIC Link is Up 1000 Mbps Full Duplex, Flow Control: RX
[ 13.791790] IPv6: ADDRCONF(NETDEV_CHANGE): enp0s3: link becomes ready
[ 14.323689] vboxvideo: loading version 6.1.38 r153438
[ 14.461936] 05:16:55.355475 main      VBoxService 6.1.38 r153438 (verbosity: 0) linux.amd64 (Sep  1 2022 15:42:08) release lo
g
05:16:55.355477 main      Log opened 2022-11-03T05:16:55.355469000Z
[ 14.461992] 05:16:55.355557 main      OS Product: Linux
[ 14.462038] 05:16:55.355604 main      OS Release: 5.4.214syslab
[ 14.462084] 05:16:55.355650 main      OS Version: #7 SMP Wed Oct 5 19:13:13 KST 2022
[ 14.462137] 05:16:55.355695 main      Executable: /opt/VBoxGuestAdditions-6.1.38/sbin/VBoxService
05:16:55.355696 main      Process ID: 918
05:16:55.355696 main      Package type: LINUX_64BITS_GENERIC
[ 14.469020] 05:16:55.362577 main      6.1.38 r153438 started. Verbose level = 0
[ 14.470993] 05:16:55.364550 main      vbglR3GuestCtrlDetectPeekGetCancelSupport: Supported (#1)
[ 14.526693] vboxsf: g_fHostFeatures=0x8000000f g_fSfFeatures=0x1 g_uSfLastFunction=29
[ 14.526752] *** VALIDATE vboxsf ***
[ 14.526756] vboxsf: Successfully loaded version 6.1.38 r153438
[ 14.526797] vboxsf: Successfully loaded version 6.1.38 r153438 on 5.4.214syslab SMP mod_unload modversions (LINUX_VERSION_C
ODE=0x504d6)
[ 14.532183] 05:16:55.425725 automount vbsvcAutomounterMountIt: Successfully mounted 'ubuntu' on '/home/syslab/ubuntu'
[ 20.473984] rfkill: input handler disabled
[ 323.001912] hrtimer: interrupt took 9520241 ns
[12496.736561] PXT4-fs: Unable to register as ext3 (-16)
[12498.312112] PXT4-fs (sdb): mounted filesystem with ordered data mode. Opts: (pull)
[12613.894777] file_write_iter is called 3,145,728 times, and the time interval is 169,827,952,250ns
[12613.894786] cpu[0] called pxt4_file_write_iter() 839701 times (26%)
[12613.894790] cpu[1] called pxt4_file_write_iter() 854459 times (27%)
[12613.894793] cpu[2] called pxt4_file_write_iter() 641532 times (20%)
[12613.894795] cpu[3] called pxt4_file_write_iter() 810036 times (25%)
syslab@syslab-vmtoolsdbox:~/pxt4$
```

# Tips

## ❖ Get the currently running task's task\_struct with macro "current":

```
static __always_inline struct task_struct *get_current(void)
{
    return this_cpu_read_stable(current_task);
}

#define current get_current()
```

From arch/x86/include/asm/current.h

## ❖ Get current task's CPU id with:

```
current->cpu
```

### ■ Cf> struct task\_struct

```
struct task_struct {
    ...
    /* Current CPU: */
    unsigned int          cpu;
    ...
}
```