

Profile Functions with Per-core Calclock

Practical Class 9

Systems and Storage Laboratory

Department of Computer Science and Engineering

Chung-Ang University

Contents

1. Why Per-core Calclock?
2. Per-core Calclock vs Original Calclock
3. Implementation

1. Why Per-core Calclock?

- Original calclock adding to kernel (calclock from practical class 7-a) can **degrade** system performance
- Per-core calclock adding to kernel **does not** much affect system performance

2. Per-core Calclock vs Original Calclock

Per-core Calclock

```
#include "calclock.h"

KTDEF(function_A);
int foo(void) {
    int ret;
    ktime_t stopwatch[2];

    ktget(&stopwatch[0]);
    ret = function_A();
    ktget(&stopwatch[1]);
    ktput(stopwatch, function_A);

    return ret;
}
```

Original Calclock

```
#include "calclock.h"

unsigned long long count, time;
int foo(void) {
    int ret;
    struct timespec stopwatch[2];

    getrawmonotonic(&stopwatch[0]);
    ret = function_A();
    getrawmonotonic(&stopwatch[1]);
    calclock(stopwatch, &time, &count);

    return ret;
}
```

Per-core Calclock	Original Calclock
KTDEF(func_name)	unsigned long long count, time
ktime_t	struct timespec
ktget	getrawmonotonic
ktput	calclock

2. Per-core Calclock vs Original Calclock

Per-core Calclock

```
#include "caltclock.h"

KTDEC(function_A);
void exit_module(void) {
    ktprint(1, function_A);
}
```

Original Calclock

```
#include "caltclock.h"

extern unsigned long long count, time;
void exit_module(void) {
    printk("Function A is called %llu times, and the time interval is %llu ns\n",
count, time);
}
```

Per-core Calclock

KTDEC(func_name)

Original Calclock

extern unsigned long long count, time

3. Implementation

- Firstly, in calclock.h, **define CONFIG_CALCLOCK**

calloc.h

```
#ifndef __CALCLOCK_H
#define __CALCLOCK_H
#include <linux/ktime.h>
#include <linux/percpu.h>

#define CONFIG_CALCLOCK

struct calclock {
    ktime_t time;
    unsigned long long count;
};

// structs and functions are defined here!

#else /* !CONFIG_CALCLOCK */
#define ktget(clock)
#define ktput(localclock, funcname)
#define ktprint(depth, funcname)
#endif /* CONFIG_CALCLOCK */

#endif /* __CALCLOCK_H */
```

3. Implementation

- Calclock structure

calloclock.h

```
struct calloclock {  
    ktime_t time;  
    unsigned long long count;  
};
```

- **time**: Time interval of the function being measured
- **count**: The number of this calloclock called by its thread

3. Implementation

- **KTDEF:** Defines a new calclock

calclock.h

```
#define KTDEF(funcname) \  
    DEFINE_PER_CPU(struct calclock, funcname##_clock) = {0, 0}
```

- **KTDEC:** Declares an existing calclock defined in another file

calclock.h

```
#define KTDEC(funcname) \  
    DECLARE_PER_CPU(struct calclock, funcname##_clock)
```


3. Implementation

- **ktget**: Gets current time and saves it into a **local clock**

❑ **Local clock** is a pointer of **ktime_t**

calclock.h

```
static inline void ktget(ktime_t *clock)
{
    *clock = ktime_get_raw();
}
```

Example:

foo.c

```
#include "calclock.h"

KTDEF(function_A);
int foo(void) {
    int ret;
    ktime_t stopwatch[2];

    ktget(&stopwatch[0]);
    ret = function_A();
    ktget(&stopwatch[1]);
    ktput(stopwatch, function_A);

    return ret;
}
```

2 local clocks of *ktime_t* are declared

ktget gets the start time and end time of 'function_A', saves them into 'stopwatch' local clocks

3. Implementation

- **ktput**: gets current per-core calclock and calculates time interval
- ❑ **localclocks**: Local clocks of ktime_t
- ❑ **funcname**: Name of per-CPU calclock (named freely)

```
calclock.h  
  
#define ktput(localclocks, funcname) \\\br/>do { \\\br/>    struct calclock *clock; \\\br/>    bool prmpt_enabled = preemptible(); \\\br/> \\\br/>    if (prmpt_enabled) \\\br/>        preempt_disable(); \\\br/>    clock = this_cpu_ptr(&(funcname##_clock)); \\\br/>    __ktput(localclocks, &clock->time); \\\br/>    clock->count++; \\\br/>    if (prmpt_enabled) \\\br/>        put_cpu_ptr(&(funcname##_clock)); \\\br/>} while (0)
```

3. Implementation

- **ktprint**: prints total time and counts of the per-core calclock
- ❑ **depth**: The number of "tab" to the left side
- ❑ **funcname**: Name of per-core calclock (named freely)

calloc.h

```
#define ktprint(depth, funcname) \
do { \
    int cpu; \
    ktime_t timesum = 0; \
    unsigned long long countsum = 0; \
\
    for_each_online_cpu(cpu) { \
        struct calclock *clock = per_cpu_ptr(&funcname##_clock, cpu); \
        timesum += clock->time; \
        countsum += clock->count; \
    } \
    __ktprint(depth, #funcname, timesum, countsum); \
} while (0)
```

3. Implementation

- Refer uploaded **calclock.h**, **calclock.c** files at (<https://github.com/loglamo/CAU-CSE-LinuxApplications-20232>)