

Assignment-10

Linux System and its Applications

Systems and Storage Laboratory

Department of Computer Science and Engineering

Chung-Ang University

Assignment-10: Synchronization atomic instruction

❖ Atomic operations (atomic instructions)

- Perform atomic operations using atomic instructions in a simple kernel module with four kernel threads.
 - Fetch-and-add
 - Test-and-set
 - Compare-and-swap
- Each thread should increase the shared resource “counter” by 1 and print it.
- Reference:
 - Lecture slide **6. synchronization (2)**
 - <https://www.ibm.com/docs/en/xcfbg/121.141?topic=functions-gcc-atomic-memory-access-built-in>

Example screenshot 1 atomic instruction

❖ Atomic operations

- Compare-and-Swap example after inserting module

```
[ 1655.812340] compare_and_swap_module_init: Entering Compare and swap Module!  
[ 1655.813086] pid[5906] compare_and_swap_function: counter: 0  
[ 1655.813330] pid[5907] compare_and_swap_function: counter: 1  
[ 1655.813527] pid[5908] compare_and_swap_function: counter: 2  
[ 1655.813711] pid[5909] compare_and_swap_function: counter: 3  
[ 1656.346667] pid[5908] compare_and_swap_function: counter: 4  
[ 1656.346682] pid[5907] compare_and_swap_function: counter: 5  
[ 1656.346688] pid[5906] compare_and_swap_function: counter: 6  
[ 1656.346716] pid[5909] compare_and_swap_function: counter: 7  
[ 1656.858630] pid[5909] compare_and_swap_function: counter: 8  
[ 1656.858632] pid[5906] compare_and_swap_function: counter: 9  
[ 1656.858633] pid[5907] compare_and_swap_function: counter: 10  
[ 1656.858635] pid[5908] compare_and_swap_function: counter: 11  
[ 1657.370292] pid[5908] compare_and_swap_function: counter: 12  
[ 1657.370298] pid[5907] compare_and_swap_function: counter: 13  
[ 1657.370301] pid[5906] compare_and_swap_function: counter: 14  
[ 1657.370304] pid[5909] compare_and_swap_function: counter: 15
```

- When removing module

```
[ 1659.932640] compare_and_swap_module_cleanup: Exiting Compare and Swap Module!
```

❖ Linked list with synchronization

- Protect linked list operations (such as insert, search, and delete) by using three different locking mechanisms in your kernel module with four kernel threads.
 - Spinlock
 - Mutex
 - RW semaphore

❖ Linked list with synchronization scenario

- Four kernel thread runs simultaneously.
- Each kernel thread perform Insert, Search, Delete operations to the global linked list, which is shared by all threads.
- To protect global linked list from corruption by concurrent access of the threads, single global lock is used.
- Each thread perform Insert, Search, Delete for its own data range bound.

❖ Linked list with synchronization

■ Tips.

- Create four threads in the module initialization procedure by using **work_fn()** on the next slide.
- Stop four threads in the module cleanup procedure by using **kthread_stop()**.
- You might need **delay.h** header file provided by Linux kernel.

Assignment-10: Synchronization basic locking

❖ Worker function example

```
static int work_fn(void *data)
{
    int range_bound[2];
    int thread_id = *(int*) data;

    set_iter_range(thread_id, range_bound);
    void *ret = add_to_list(thread_id, range_bound);
    search_list(thread_id, ret, range_bound);
    delete_from_list(thread_id, range_bound);

    while(!kthread_should_stop()) {
        msleep(500);
    }
    printk(KERN_INFO "thread #%d stopped!\n", thread_id);

    return 0;
}
```

Assignment-10: Synchronization basic locking

❖ Linked list with synchronization example

- You have to **fill in the blank** and complete the code **linked_list_impl.c**
- Different implementation is okay for the submission.

```
#include "calclock.h"
```

```
// define your spinlock here
```

```
// initialize your list here
```

```
void *add_to_list(int thread_id, int range_bound[])  
{
```

```
...
```

```
printk(KERN_INFO "thread #%d range: %d ~ %d\n",  
        thread_id, range_bound[0], range_bound[1]);
```

```
// put your code here
```

```
return first;
```

```
// return first entry that was inserted with the current thread
```

```
}
```

```
int search_list(int thread_id, void *data, int range_bound[])  
{
```

```
struct timespec localclock[2];
```

```
/* This will point on the actual data structures during the iteration */
```

```
struct my_node *cur = (struct my_node *) data, *tmp;
```

```
// put your code here
```

```
return 0;
```

```
}
```

```
int delete_from_list(int thread_id, int range_bound[])  
{
```

```
struct my_node *cur, *tmp;
```

```
struct timespec localclock[2];
```

```
// put your code here
```

```
return 0;
```

```
}
```


Example screenshot 1 basic locking

❖ Linked list with synchronization

- Please follow the below template when printing out
- Insert, Search, Delete example after inserting module

```
[ 3922.947470] spinlock_module_init: Entering Spinlock Module!  
[ 3922.948019] thread #1 range: 0 ~ 249999  
[ 3922.948380] thread #2 range: 250000 ~ 499999  
[ 3922.948750] thread #3 range: 500000 ~ 749999  
[ 3922.953777] thread #4 range: 750000 ~ 999999  
[ 3923.010979] thread #2 searched range: 250000 ~ 499999  
[ 3923.048050] thread #3 searched range: 500000 ~ 749999  
[ 3923.084652] thread #4 searched range: 750000 ~ 999999  
[ 3923.098999] thread #1 searched range: 0 ~ 249999  
[ 3923.114864] thread #1 deleted range: 0 ~ 249999  
[ 3923.131260] thread #2 deleted range: 250000 ~ 499999  
[ 3923.148294] thread #3 deleted range: 500000 ~ 749999  
[ 3923.165111] thread #4 deleted range: 750000 ~ 999999
```

- When removing module

```
[ 4299.721428] spinlock_module_cleanup: Spinlock linked list insert time: 59072192 ns, count: 1000000  
[ 4299.721429] spinlock_module_cleanup: Spinlock linked list search time: 20077889 ns, count: 1000000  
[ 4299.721430] spinlock_module_cleanup: Spinlock linked list delete time: 39750986 ns, count: 1000000  
[ 4300.162326] thread #1 stopped!  
[ 4300.190519] thread #2 stopped!  
[ 4300.222132] thread #3 stopped!  
[ 4300.253640] thread #4 stopped!  
[ 4300.253923] spinlock_module_cleanup: Exiting Spinlock Module!
```

What to submit

❖ Atomic operations

- Short summary of each atomic operation.
- Code screenshot of each atomic operation module file.
- Screenshot of dmesg after successful insertion & removal of module for each atomic operation.

❖ Linked list with synchronization

- Code screenshot of each locking mechanism module's `linked_list_impl.c` file.
- Screenshot of time measure result of each operation (insert, search, delete) at 1,000,000 nodes (250,000 per thread) while using spinlock, mutex, RW semaphore each.

What to submit

❖ Submission format should be pdf.

- Make sure to include `linked_list_impl.c` screenshot and **dmesg** screenshots in submitted PDF.
- Make sure to include your name and student id.